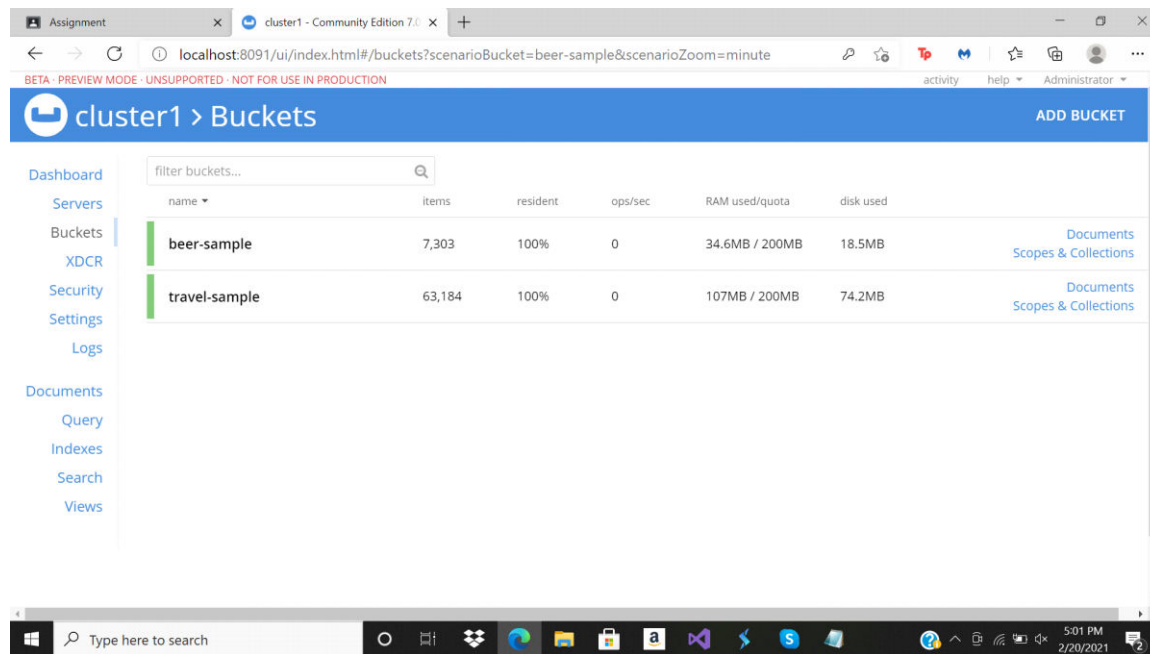
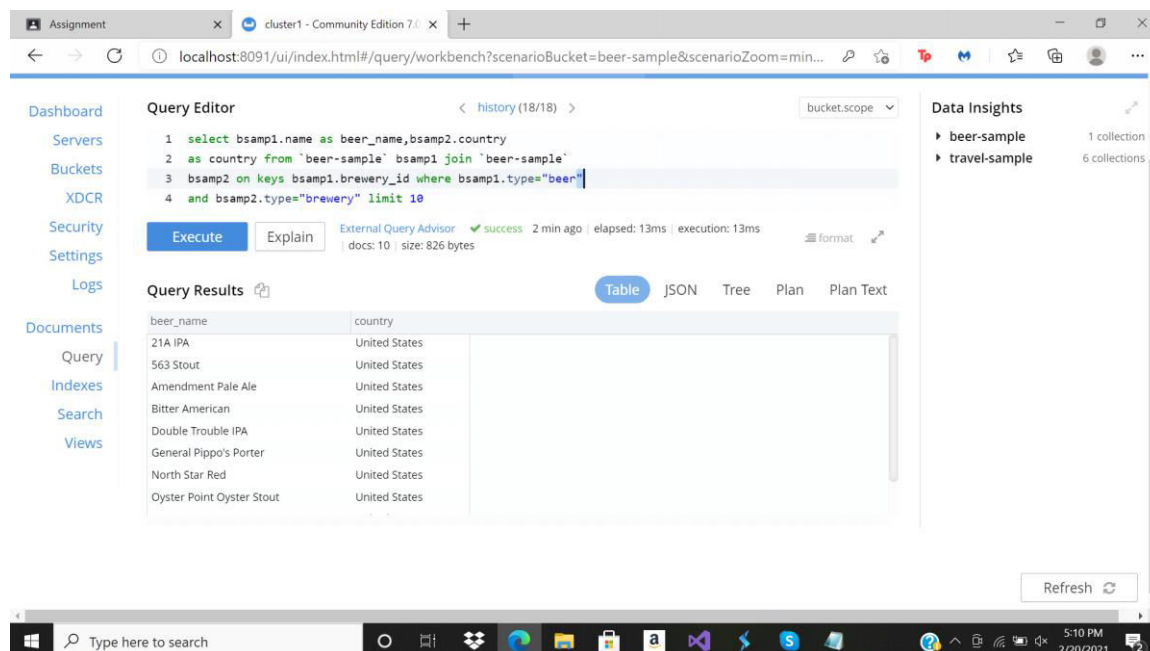


1. Import `beer-sample` bucket



2. Write a join query to fetch the Top 10 brewery(type="beer") and their country(type="brewery") which produces more varieties of beers.



3. Write a mapreduce to get the number of breweries based on country.
Please attach the mapreduce code and json output screenshot.

The screenshot shows a web browser window with a URL bar indicating a local development environment. The page displays a MapReduce job configuration and its results.

Map Code:

```
1 function (doc, meta) {  
2   if(doc.type=="brewery")  
3     emit(doc.country,null);  
4 }
```

Reduce Code (built in: _count, _sum, _stats):

```
1
```

Results: filter: ?limit=6&stale=false&connection_timeout=60000&inclusive_end=true&skip=0&full_set=

Development Time Subset **Full Cluster Data Set** **Show Results**

Key	Value
"21st_amendment_brewery_cafe" 21st_amendment_brewery_cafe	null
"21st_amendment_brewery_cafe-21a_ipa" 21st_amendment_brewery_cafe-21a_ipa	null
"21st_amendment_brewery_cafe-563_stout" 21st_amendment_brewery_cafe-563_stout	null
"21st_amendment_brewery_cafe-amendment_pale_ale" 21st_amendment_brewery_cafe-amendment_pale_ale	null
"21st_amendment_brewery_cafe-bitter_american" 21st_amendment_brewery_cafe-bitter_american	null
"21st_amendment_brewery_cafe-double_trouble_ipa" 21st_amendment_brewery_cafe-double_trouble_ipa	null

4. XDCR:

a) Add a new bucket “Brewery”.

b) Create a XDCR with a filter(type='brewery') to replicate only the brewery entity from `beer-sample` bucket.

The top screenshot shows the 'cluster1 > Buckets' page. It features a sidebar with navigation links: Dashboard, Servers, Buckets, XDCR, Security, Settings, Logs, Documents, Query, Indexes, Search, and Views. The main content area has a search bar 'filter buckets...' and a table of buckets.

name	items	resident	ops/sec	RAM used/quota	disk used	
beer-sample	7,303	100%	0	37.9MB / 200MB	32.9MB	Documents Scopes & Collections
brewery	7,303	100%	0	32.8MB / 4.14GB	9.58MB	Documents Scopes & Collections
travel-sample	63,184	100%	0	107MB / 200MB	74.2MB	Documents Scopes & Collections

The bottom screenshot shows the 'cluster1 > XDCR Replications' page. It has a sidebar with the same navigation links. The main content area has a search bar and two sections: 'Remote Clusters' and 'Outgoing Replications'.

Remote Clusters

name	IP/hostname
buc1	127.0.0.1:8091

Outgoing Replications

source bucket	destination bucket	remote cluster	status
beer-sample	brewery	buc1	replicating ..

